

Compiling and fitting a model

Puteaux, Fall/Winter 2020-2021

```
#####  
##                               ##  
## Deep Learning in Python      ##  
##                               ##  
#####
```

§1 Introduction to Deep Learning in Python

§1.3 Building deep learning models with keras

§1.3.2 Compiling and fitting a model

1. Why is it necessary to compile the model?

- Specify the optimizer:
 - many options and mathematically complex
 - adam is usually a good choice
- Loss function:
 - mean_squared_error is common for regression

2. Code of compiling a model:

```
[1]: import numpy as np  
from keras.layers import Dense  
from keras.models import Sequential  
  
predictors = np.loadtxt('ref1. Hourly wages predictors data.csv',  
                        delimiter=',')  
n_cols = predictors.shape[1]  
  
model = Sequential()  
model.add(Dense(100, activation='relu', input_shape=(n_cols, )))  
model.add(Dense(100, activation='relu'))  
model.add(Dense(1))  
model.compile(optimizer='adam', loss='mean_squared_error')
```

3. What is fitting a model?

- Apply backpropagation and gradient descent with the data to update the weights.

- Scale data before fitting can ease optimization.

4. Code of fitting a model:

```
[2]: target = np.loadtxt('ref3. Hourly wages target data.csv', delimiter=',')

model.fit(predictors, target)
```

17/17 [=====] - 1s 41ms/step - loss: 422.6907

```
[2]: <tensorflow.python.keras.callbacks.History at 0x7fa93625abd0>
```

5. Practice exercises for compiling and fitting a model:

► Package pre-loading:

```
[3]: import pandas as pd
```

► Data pre-loading:

```
[4]: df = pd.read_csv('ref2. Hourly wages.csv')

predictors = df.iloc[:, 1:].to_numpy()
```

► Model compiling practice:

```
[5]: # Import necessary modules
import keras
from keras.layers import Dense
from keras.models import Sequential

# Specify the model
n_cols = predictors.shape[1]
model = Sequential()
model.add(Dense(50, activation='relu', input_shape=(n_cols, )))
model.add(Dense(32, activation='relu'))
model.add(Dense(1))

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Verify that model contains information from compiling
print("Loss function: " + model.loss)
```

Loss function: mean_squared_error

► Data re-pre-loading:

```
[6]: target = df.iloc[:, 0].to_numpy()
```

► Model fitting practice:

```
[7]: # Import necessary modules
import keras
from keras.layers import Dense
from keras.models import Sequential

# Specify the model
n_cols = predictors.shape[1]
model = Sequential()
model.add(Dense(50, activation='relu', input_shape=(n_cols, )))
model.add(Dense(32, activation='relu'))
model.add(Dense(1))

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Fit the model
model.fit(predictors, target)
```

17/17 [=====] - 1s 59ms/step - loss: 32.0553

[7]: <tensorflow.python.keras.callbacks.History at 0x7fa9368a0590>