

Loss functions

Puteaux, Fall/Winter 2020-2021

```
#####  
##                               ##  
## Deep Learning in Python  ##  
##                               ##  
#####
```

§2 Introduction to TensorFlow in Python

§2.2 Linear models

1 Loss functions

1.1 What is the loss functions?


- **Fundamental TensorFlow operation:**
 - used to train a model
 - measure of model fit
- **Higher value → worse fit:**
 - minimize the loss function

1.2 What are the common loss functions in TensorFlow?

- **TensorFlow has operations for common loss functions:**
 - mean squared error (MSE)
 - mean absolute error (MAE)
 - huber error
- **Loss functions are accessible from `tf.keras.losses()`:**
 - `tf.keras.losses.mse()`
 - `tf.keras.losses.mae()`
 - `tf.keras.losses.Huber()`

1.3 Why is it in need to care about loss functions?

- **MSE:**
 - strongly penalizes outliers
 - high (gradient) sensitivity near minimum
- **MAE:**
 - scales linearly with the size of the error
 - low sensitivity near minimum
- **Huber:**
 - similar to MSE near minimum
 - similar to MAE away from the minimum



1.4 Code of defining the loss function:

```
[1]: import pandas as pd
import numpy as np

housing = pd.read_csv('ref1. King county house sales.csv')

price_log = np.log(np.array(housing['price'], np.float32))
size_log = np.log(np.array(housing['sqft_lot'], np.float32))
targets = price_log
features = size_log
intercept = 0.1
slope = 0.1

predictions = intercept + features * slope
```

```
[2]: # Import TensorFlow under standard alias
import tensorflow as tf

# Compute the MSE loss
loss = tf.keras.losses.mse(targets, predictions)

loss
```

```
[2]: <tf.Tensor: shape=(), dtype=float32, numpy=145.44653>
```

```
[3]: # Define a linear regression model
def linear_regression(intercept, slope=slope, features=features):
    return intercept + features * slope
```

```
[4]: # Define a loss function to compute the MSE
def loss_function(intercept, slope, targets=targets, features=features):
    # Compute the predictions for a linear model
    predictions = linear_regression(intercept, slope, features)

    # Return the loss
    return tf.keras.losses.mse(targets, predictions)
```

```
[5]: from sklearn.model_selection import train_test_split

train_features, test_features, train_targets, test_targets = train_test_split(
    features, targets, test_size=0.3, random_state=42)
```

```
[6]: # Compute the loss for test data inputs
loss_function(intercept, slope, test_targets, test_features)
```

```
[6]: <tf.Tensor: shape=(), dtype=float32, numpy=145.57338>
```

```
[7]: # Compute the loss for default data inputs
loss_function(intercept, slope)
```

```
[7]: <tf.Tensor: shape=(), dtype=float32, numpy=145.44653>
```

1.5 Practice exercises for loss functions:

► Package pre-loading:

```
[8]: import pandas as pd
import numpy as np
```

► Data pre-loading:

```
[9]: housing = pd.read_csv('ref1. King county house sales.csv')
price = np.array(housing['price'], np.float32)

predictions = np.loadtxt('ref5. Predictions.csv', delimiter=',')
```

► TensorFlow loss functions practice:

```
[10]: # Import the keras module from tensorflow
from tensorflow import keras

# Compute the mean squared error (mse)
loss = keras.losses.mse(price, predictions)

# Print the mean squared error (mse)
print(loss.numpy())
```

141171604777.141

```
[11]: # Import the keras module from tensorflow
from tensorflow import keras

# Compute the mean absolute error (mae)
loss = keras.losses.mae(price, predictions)

# Print the mean absolute error (mae)
print(loss.numpy())
```

268827.9930163703

► Package re-pre-loading:

```
[12]: from tensorflow import constant, Variable, float32
```

► Data re-pre-loading:

```
[13]: features = constant([1., 2., 3., 4., 5.], dtype=float32)
targets = constant([2., 4., 6., 8., 10.], dtype=float32)
```

► Loss function modification practice:

```
[14]: # Initialize a variable named scalar
scalar = Variable(1.0, float32)

# Define the model
def model(scalar, features=features):
    return scalar * features
```

```
# Define a loss function
def loss_function(scalar, features=features, targets=targets):
    # Compute the predicted values
    predictions = model(scalar, features)

    # Return the mean absolute error loss
    return keras.losses.mae(targets, predictions)

# Evaluate the loss function and print the loss
print(loss_function(scalar).numpy())
```

3.0

