

Introduction to gensim

Puteaux, Fall/Winter 2020-2021

```
#####  
##                                     ##  
## Natural Language Processing in Python ##  
##                                     ##  
#####
```

§1 Introduction to Natural Language Processing in Python


§1.2 Simple topic identification

1 Introduction to gensim

1.1 What is gensim?

- It is a popular open-source NLP library.
- It uses top academic models to perform complex tasks:
 - building document or word vectors
 - performing topic identification and document comparison

1.2 What is a word vector?



1.3 Code of creating a gensim corpus:

```
[1]: from gensim.corpora.dictionary import Dictionary  
from nltk.tokenize import word_tokenize  
  
my_documents = [  
    'The movie was about a spaceship and aliens.',
```

```
'I really liked the movie!',  
'Awesome action scenes, but boring characters.',  
'The movie was awful! I hate alien films.',  
'Space is cool! I liked the movie.',  
'More space films, please!'],
```

```
[2]: tokenized_docs = [word_tokenize(doc.lower()) for doc in my_documents]  
  
dictionary = Dictionary(tokenized_docs)  
dictionary.token2id
```

```
[2]: {'.': 0,  
      'a': 1,  
      'about': 2,  
      'aliens': 3,  
      'and': 4,  
      'movie': 5,  
      'spaceship': 6,  
      'the': 7,  
      'was': 8,  
      '!': 9,  
      'i': 10,  
      'liked': 11,  
      'really': 12,  
      ',': 13,  
      'action': 14,  
      'awesome': 15,  
      'boring': 16,  
      'but': 17,  
      'characters': 18,  
      'scenes': 19,  
      'alien': 20,  
      'awful': 21,  
      'films': 22,  
      'hate': 23,  
      'cool': 24,  
      'is': 25,  
      'space': 26,  
      'more': 27,  
      'please': 28}
```

```
[3]: corpus = [dictionary.doc2bow(doc) for doc in tokenized_docs]  
corpus
```

```
[3]: [[(0, 1), (1, 1), (2, 1), (3, 1), (4, 1), (5, 1), (6, 1), (7, 1), (8, 1)],  
      [(5, 1), (7, 1), (9, 1), (10, 1), (11, 1), (12, 1)],
```

```
[(0, 1), (13, 1), (14, 1), (15, 1), (16, 1), (17, 1), (18, 1), (19, 1)],  
[(0, 1),  
 (5, 1),  
 (7, 1),  
 (8, 1),  
 (9, 1),  
 (10, 1),  
 (20, 1),  
 (21, 1),  
 (22, 1),  
 (23, 1)],  
[(0, 1), (5, 1), (7, 1), (9, 1), (10, 1), (11, 1), (24, 1), (25, 1), (26, 1)],  
[(9, 1), (13, 1), (22, 1), (26, 1), (27, 1), (28, 1)]]
```

1.4 What are the advantages of creating a gensim corpus?

- First of all, gensim models can be easily saved, updated, and reused.
- Secondly, the dictionary created can also be updated.
- Lastly, the more advanced and feature-rich bag-of-words can be used in future exercises.

1.5 Practice question for word vectors:

- What are word vectors, and how do they help with NLP?
 - ☐ They are similar to bags of words, just with numbers. You use them to count how many tokens there are.
 - ☐ Word vectors are sparse arrays representing bigrams in the corpora. You can use them to compare two sets of words to one another.
 - ☒ Word vectors are multi-dimensional mathematical representations of words created using deep learning methods. They give us insight into relationships between words in a corpus.
 - ☐ Word vectors don't actually help NLP and are just hype.

1.6 Practice exercises for introduction to gensim:

► Package pre-loading:

```
[4]: import zipfile  
  
from nltk import word_tokenize  
from nltk.corpus import stopwords
```

► Data pre-loading:

```
[5]: file_name = 'ref3. Wikipedia articles.zip'  
with zipfile.ZipFile(file_name, 'r') as archive:  
    files = [
```

```
        archive.read(name) for name in archive.namelist()
        if name.endswith('.txt')
    ]

doc_tokens = [word_tokenize(file.decode("utf-8")) for file in files]

articles = []
english_stops = stopwords.words('english')
for i in range(len(doc_tokens)):
    lower_tokens = [t.lower() for t in doc_tokens[i]]
    alphanumeric_only = [t for t in lower_tokens if t.isalnum()]
    no_stops = [t for t in alphanumeric_only if t not in english_stops]
    articles.append(no_stops)
```

► Gensim corpus creating and querying practice:

```
[6]: # Import Dictionary
from gensim.corpora.dictionary import Dictionary

# Create a Dictionary from the articles: dictionary
dictionary = Dictionary(articles)

# Select the id for "computer": computer_id
computer_id = dictionary.token2id.get("computer")

# Use computer_id with the dictionary to print the word
print(dictionary.get(computer_id))

# Create a MmCorpus: corpus
corpus = [dictionary.doc2bow(article) for article in articles]

# Print the first 10 word ids with their frequency counts from the fifth
# document
print(corpus[4][:10])
```

```
computer
[(13, 2), (24, 1), (43, 1), (44, 6), (45, 1), (50, 1), (58, 1), (59, 1), (61,
7), (75, 1)]
```

► Package pre-loading:

```
[7]: from collections import defaultdict
import itertools
```

► Gensim bag-of-words practice:

```
[8]: # Save the fifth document: doc
doc = corpus[4]
```

```
# Sort the doc for frequency: bow_doc
bow_doc = sorted(doc, key=lambda w: w[1], reverse=True)

# Print the top 5 words of the document alongside the count
for word_id, word_count in bow_doc[:5]:
    print(dictionary.get(word_id), word_count)

# Create the defaultdict: total_word_count
total_word_count = defaultdict(int)
for word_id, word_count in itertools.chain.from_iterable(corpus):
    total_word_count[word_id] += word_count
```

language 54
programming 39
languages 30
code 22
computer 15

```
[9]: # Save the fifth document: doc
doc = corpus[4]

# Sort the doc for frequency: bow_doc
bow_doc = sorted(doc, key=lambda w: w[1], reverse=True)

# Print the top 5 words of the document alongside the count
for word_id, word_count in bow_doc[:5]:
    print(dictionary.get(word_id), word_count)

# Create the defaultdict: total_word_count
total_word_count = defaultdict(int)
for word_id, word_count in itertools.chain.from_iterable(corpus):
    total_word_count[word_id] += word_count

# Create a sorted list from the defaultdict: sorted_word_count
sorted_word_count = sorted(total_word_count.items(),
                           key=lambda w: w[1],
                           reverse=True)

# Print the top 5 words across all documents alongside the count
for word_id, word_count in sorted_word_count[:5]:
    print(dictionary.get(word_id), word_count)
```

language 54
programming 39
languages 30
code 22
computer 15

computer 598

software 450

cite 322

ref 259

code 235

